

# HW5 - Linear & Logistic regression

February 18, 2019

## 1 Data-X Spring 2019: Homework 05

### 1.0.1 Linear regression & Logistic regression

**1.1 Name: Shrey Samdani**

**1.2 SID: 3032000414**

In this homework, you will do some exercises on prediction using sklearn.

REMEMBER TO DISPLAY ALL OUTPUTS. If the question asks you to do something, make sure to print your results.

### 1.3 Part 1 - Regression

#### 1.3.1 Data:

**Data Source:** Data file is uploaded to bCourses and is named: **Energy.csv** (Link in the Assignment details page on Bcourses)

The dataset was created by Angeliki Xifara ( Civil/Structural Engineer) and was processed by Athanasios Tsanas, Oxford Centre for Industrial and Applied Mathematics, University of Oxford, UK).

#### **Data Description:**

The dataset contains eight attributes of a building (or features, denoted by X1...X8) and response being the heating load on the building, y1.

- X1 Relative Compactness
- X2 Surface Area
- X3 Wall Area
- X4 Roof Area
- X5 Overall Height
- X6 Orientation
- X7 Glazing Area
- X8 Glazing Area Distribution
- y1 Heating Load

#### **Q1.1** Read the data file from the csv.

Print the count of NaN values for each attribute in the dataset.

Print the Range (min, max) and percentiles (25th, 50th, and 75th) of each attribute in the dataset

```
In [11]: import numpy as np
import pandas as pd

energy = pd.read_csv("energy.csv")
print("# of NaN values:", energy.isnull().sum().sum())
energy.describe()
```

# of NaN values: 0

```
Out[11]:
```

	X1	X2	X3	X4	X5	X6 \
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	0.764167	671.708333	318.500000	176.604167	5.250000	3.500000
std	0.105777	88.086116	43.626481	45.165950	1.75114	1.118763
min	0.620000	514.500000	245.000000	110.250000	3.500000	2.000000
25%	0.682500	606.375000	294.000000	140.875000	3.500000	2.750000
50%	0.750000	673.750000	318.500000	183.750000	5.250000	3.500000
75%	0.830000	741.125000	343.000000	220.500000	7.000000	4.250000
max	0.980000	808.500000	416.500000	220.500000	7.000000	5.000000

  

	X7	X8	Y1
count	768.000000	768.000000	768.000000
mean	0.234375	2.81250	22.307201
std	0.133221	1.55096	10.090196
min	0.000000	0.00000	6.010000
25%	0.100000	1.75000	12.992500
50%	0.250000	3.00000	18.950000
75%	0.400000	4.00000	31.667500
max	0.400000	5.00000	43.100000

## REGRESSION:

Using the data, we want to predict "Heating load". The output variable is continuous. Hence, we need to use a regression algorithm.

### Q 1.2:

Split the dataset randomly into train and test. Train a **Linear Regression** model on 80% of the data (80-20 split). What is the intercept and coefficient values?

```
In [12]: from sklearn.model_selection import train_test_split

X = energy.iloc[:, :-1]
Y = energy['Y1']
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)

from sklearn import linear_model

linearModel = linear_model.LinearRegression()
linearModel.fit(x_train, y_train)
```

```
print("Intercept:",linearModel.intercept_)
print("Coefficients:", linearModel.coef_)
```

Intercept: 73.63857288273616

Coefficients: [-5.79848476e+01 -1.39460642e+11 1.39460642e+11 2.78921284e+11  
4.13473804e+00 -2.13875247e-02 1.98411253e+01 2.32395440e-01]

**Q.1.3:** Create a function which takes arrays of prediction and actual values of the output as parameters to calculate 'Root Mean Square error' (RMSE) metric:

1. Use the function to calculate the training RMSE
2. Use the function to calculate the test RMSE

```
In [13]: # your code
rmseCalc = lambda pred, actual: np.mean((pred-actual)**2)**0.5

print("TRAINING RMSE:", rmseCalc(linearModel.predict(x_train), y_train))
print("TEST RMSE:", rmseCalc(linearModel.predict(x_test), y_test))
```

TRAINING RMSE: 2.892919795496745

TEST RMSE: 3.03337797694144

**Q1.4:** Let's see the effect of amount of data on the performance of prediction model. Use varying amounts of data (100,200,300,400,500,all) from the training data you used previously to train different regression models. Report training error and test error in each case. Test data is the same as above for all these cases.

**Plot error rates vs number of training examples.** Both the training error and the test error should be plotted. Comment on the relationship you observe between the amount of data used to train the model and the test accuracy of the model.

**Hint:** Use array indexing to choose varying data amounts

```
In [14]: import matplotlib.pyplot as plt

train_sizes = [100,200,300,400,500, len(x_train)]
train_error, test_error = [], []

for size in train_sizes:
    linearModel = linear_model.LinearRegression()
    linearModel.fit(x_train.iloc[:size,:], y_train.iloc[:size])
    train_error.append(linearModel.score(x_train.iloc[:size,:], y_train.iloc[:size]))
    test_error.append(linearModel.score(x_test.iloc[:size,:], y_test.iloc[:size]))

print("SIZE:",size)
print("TRAINING ACCURACY:",train_error[-1])
print("TEST ACCURACY:",test_error[-1])
```

```

print()

plt.subplot(1, 2, 1)
plt.plot(train_sizes, train_error)
plt.title("TRAINING ERROR")
plt.xlabel("SIZE")
plt.ylabel("ACCURACY")

plt.subplot(1, 2, 2)
plt.plot(train_sizes, test_error)
plt.title("TEST ERROR")
plt.xlabel("SIZE")
plt.ylabel("ACCURACY")

plt.show()

```

```

SIZE: 100
TRAINING ACCURACY: 0.9244997008930089
TEST ACCURACY: 0.9131949482613633

```

```

SIZE: 200
TRAINING ACCURACY: 0.9113355114015416
TEST ACCURACY: 0.9069831120301527

```

```

SIZE: 300
TRAINING ACCURACY: 0.922842847741354
TEST ACCURACY: 0.9067987723998217

```

```

SIZE: 400
TRAINING ACCURACY: 0.9225682273925364
TEST ACCURACY: 0.9066138628588658

```

```

SIZE: 500
TRAINING ACCURACY: 0.9214084886642311
TEST ACCURACY: 0.9064240862216149

```

```

SIZE: 614
TRAINING ACCURACY: 0.9184933052813427
TEST ACCURACY: 0.9057492207192306

```



It seems as the size increases, the accuracy goes down after a certain point, which could be explained by overfitting.

## 1.4 Part 2 - Classification

**CLASSIFICATION:** LABELS ARE DISCRETE VALUES.

Here the model is trained to classify each instance into a set of predefined discrete classes. On inputting a feature vector into the model, the trained model is able to predict a class of that instance.

**Q2.1** Bucket the values of 'y1' i.e 'Heating Load' from the original dataset into 3 classes:

- 0: 'Low' (< 14),
- 1: 'Medium' (14-28),
- 2: 'High' (>28)

**HINT:** Use pandas.cut

This converts the given dataset into a classification problem. Use this dataset with transformed 'heating load' to create a **logistic regression** classification model that predicts heating load type of a building. Split the data randomly into training and test set. Train the model on 80% of the data (80-20 split).

```
In [5]: bins = pd.IntervalIndex.from_tuples([(0, 1), (2, 3), (4, 5)])
        energy['Y_class'] = energy['Y1']
        Y = energy['Y_class']
        Y[Y < 14] = 0
```

```
Y[(Y >= 14) & (Y<=28)] = 1
Y[Y > 28] = 2
```

```
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2)
linearModel = linear_model.LogisticRegression()
linearModel.fit(x_train, y_train)
```

/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:4: SettingWithOut  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>  
after removing the cwd from sys.path.

/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:5: SettingWithOut  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>  
"""

/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:6: SettingWithOut  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear\_model/logistic.py:433  
FutureWarning)  
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear\_model/logistic.py:460  
"this warning.", FutureWarning)

```
Out[5]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, max_iter=100, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=None, solver='warn',
    tol=0.0001, verbose=0, warm_start=False)
```

## Q2.2

- Print the training and test accuracies
- Print the confusion matrix
- Print the precision and recall numbers for all the classes

```
In [6]: # your code
print("TRAINING ACCURACY:", linearModel.score(x_train, y_train))
print("TEST ACCURACY:", linearModel.score(x_test, y_test))

from sklearn.metrics import *
predicted = linearModel.predict(x_test)
print("CONFUSION MATRIX:", confusion_matrix(predicted, y_test))
print("PRECISION NUMBERS:", precision_score(predicted, y_test, average = None))
print("RECALL NUMBERS:", recall_score(predicted, y_test, average = None))
```

```

TRAINING ACCURACY: 0.7947882736156352
TEST ACCURACY: 0.7467532467532467
CONFUSION MATRIX: [[44 18  0]
 [ 2 24  0]
 [ 0 19 47]]
PRECISION NUMBERS: [0.95652174 0.39344262 1.          ]
RECALL NUMBERS: [0.70967742 0.92307692 0.71212121]

```

### Q2.3

**K Fold Cross Validation** In k-fold cross-validation, the shuffled training data is partitioned into k disjoint sets and the model is trained on k-1 sets and validated on the kth set. This process is repeated k times with each set chosen as the validation set once. The cross-validation accuracy is reported as the average accuracy of the k iterations

**Use 7-fold cross validation on the training data. Print the average accuracy**

```

In [7]: # your code
        from sklearn.model_selection import cross_val_score

        linearModel = linear_model.LogisticRegression()
        print("AVERAGE ACCURACY:", cross_val_score(linearModel, x_train, y_train).mean())

```

```
AVERAGE ACCURACY: 0.7636857275837573
```

```

/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/model_selection/_split.py:200:
  warnings.warn(CV_WARNING, FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433:
  FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:460:
  "this warning.", FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433:
  FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:460:
  "this warning.", FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433:
  FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:460:
  "this warning.", FutureWarning)

```

### Q2.4

One of the preprocessing steps in Data science is Feature Scaling i.e getting all our data on the same scale by setting same Min-Max of feature values. This makes training less sensitive to the scale of features. Scaling is important in algorithms that use distance functions as a part of classification. If we Scale features in the range [0,1] it is called unity based normalization.

**Perform unity based normalization on the above dataset and train the model again, compare model performance in training and validation with your previous model.**

refer:<http://scikit-learn.org/stable/modules/preprocessing.html#preprocessing-scaler>  
more at: [https://en.wikipedia.org/wiki/Feature\\_scaling](https://en.wikipedia.org/wiki/Feature_scaling)

```
In [8]: from sklearn.preprocessing import MinMaxScaler
        scaler = MinMaxScaler()
        scaler.fit(X)

        x_train, x_test, y_train, y_test = train_test_split(scaler.transform(X), Y, test_size=0.2)
        linearModel = linear_model.LogisticRegression()
        linearModel.fit(x_train, y_train)

        print("TRAINING ACCURACY:", linearModel.score(x_train, y_train))
        print("TEST ACCURACY:", linearModel.score(x_test, y_test))
```

```
TRAINING ACCURACY: 0.8241042345276873
TEST ACCURACY: 0.8181818181818182
```

```
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:323: DataConversionWarning:
    Data has been converted to float.
    return self.partial_fit(X, y)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning:
    The default value of 'solver' will change from 'lbfgs' to 'saga' in version 0.24.
    FutureWarning)
/Users/shreysamdani/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:460: FutureWarning:
    'this warning.', FutureWarning)
```

The model performance is very similar in training and validation with the previous model.

```
In [ ]:
```